Quantum Walk Search for Marked Vertices as a Model of Cognitive Search:

From Semantic Memory Retrieval to Multi-Alternative Decision Making

[Student Name 1] [Student ID 1]
[Student Name 2] [Student ID 2]

EE5105 Special Topics in Quantum Information Science
Instructor: Prof. Shih-Han Hung Fall 2026

Abstract

We use the Ambainis–Gilyén–Jeffery–Kokainis (AGJK) marked-vertex search algorithm as the technical core of a process-level model of cognitive search. Numerically, the quantum search time scales as $T_q \sim \mathrm{HT}^{0.46\text{–}0.54}$ on complete graphs, cycles, grids, and hypercubes, and independent targets are easier than clustered ones. A large extended-hitting-time gap appears only in the AGJK adversarial torus construction, not for generic clustered targets. On a toy semantic graph, continuous-time quantum walks produce larger and later target activation than classical diffusion, especially for cross-community targets. In a multi-alternative decision graph, coupling two options shifts the quantum choice probability toward the uncoupled option and produces non-monotonic choice probabilities for asymmetric branches. A partial- measurement sweep finds an optimal measurement rate $p \approx 0.2$ at a fixed deadline, with up to a $12\times$ improvement in deadline success, while the unrestricted peak success probability remains highest for the pure quantum walk. The project connects quantum-walk complexity to testable model-level predictions without claiming a biological quantum advantage.

1 Introduction

Quantum cognition models judgment, memory, and decision making with the mathematical structure of quantum probability. It does not claim that the brain is a quantum computer. Instead, it uses Hilbert spaces, superposition, and interference to capture contextuality, order effects, and evidence accumulation that classical probability handles awkwardly (Pothos and Busemeyer, 2022; Busemeyer and Bruza, 2024). Quantum walks are a natural process-level ingredient: a walker on a semantic or decision graph represents unfolding evidence, a marked vertex represents a target memory or an option, and the walk's statistics define response times and choice probabilities (Chen et al., 2022; Huang and Busemeyer, 2022; Pavone and Faro, 2025).

The missing link is algorithmic structure. The quantum-walk literature on search focuses on worst-case query and time complexity: given a graph and a marked set, how quickly can one detect or find a marked vertex? The quantum cognition literature focuses on explaining behavioral signatures. The two bodies of work meet mostly through the metaphor that a quantum walk is an evidence accumulator. The marked-vertex search problem itself—especially its behavior with multiple marked vertices—is rarely used as a formal model of cognitive search.

This project uses the Ambainis–Gilyén–Jeffery–Kokainis (AGJK) algorithm as the technical core (Ambainis et al., 2019). AGJK gives a quantum walk algorithm that finds a marked vertex

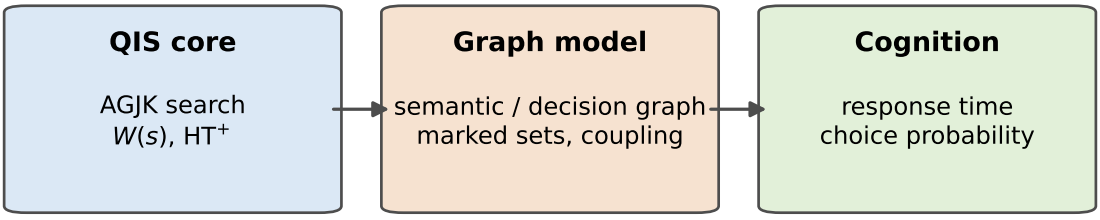


Figure 1: Three-layer architecture. The AGJK search core supplies the walk operator and the hitting-time structure; the graph layer maps concepts or options to vertices and marked sets; the cognitive layer turns walk statistics into response times and choice probabilities.

on an arbitrary graph, for an arbitrary marked set, with a quadratic speedup over the classical hitting time up to logarithmic factors. It also isolates the extended hitting time HT^+, which can be much larger than HT when there are multiple marked vertices. That structure maps naturally onto cognitive search with several targets and onto multi-alternative choice, where similar options compete and context changes the choice distribution.

We ask three questions.

1. **Mapping.** How do HT, HT^+, the spectral gap, and the success probability map onto response time, choice probability, and deliberation speed on a semantic or decision graph?

2. **Structure.** How do the number and arrangement of marked vertices—single, independent, clustered—change the quantum versus classical search time and success probability? Can HT^+ formalize multi-alternative decision difficulty and context effects?

3. **Noise.** How does partial measurement or decoherence erode the quantum advantage, and is there an optimal "introspection" level that maximizes success probability or decision speed?

Figure 1 summarizes the three layers of the project: the QIS core, the graph model, and the cognitive quantities derived from the walk.

The project makes four contributions. First, we implement the Szegedy walk, the interpolated walk $W(s)$, and the AGJK extended hitting time, and we verify the $\widetilde{O}(\sqrt{\mathrm{HT}})$ scaling numerically on complete graphs, cycles, grids, and hypercubes. Second, we compare single, independent, and clustered marked sets. We show that independent targets are easier, that clustered targets behave like a single target, and that the dramatic $\mathrm{HT}^+ \gg \mathrm{HT}$ gap requires the adversarial periodic arrangement of AGJK rather than clustering alone. Third, we build a semantic graph and a multi-branch decision graph and turn the walk statistics into target activation, response time, and choice probability. Fourth, we add a partial-measurement model and find a speed–accuracy tradeoff: weak decoherence improves deadline success and can dramatically shorten the time to half-maximum, while the unrestricted peak success probability remains highest for the pure quantum walk.

The paper is organized as follows. Section 2 reviews classical and quantum walks and the quantum cognition literature. Section 3 collects the formal definitions and the AGJK results we use. Section 4 states the cognitive mapping and its assumptions. Section 5 describes the graph families, protocols, and metrics. Section 6 reports E1–E5. Section 7 discusses interpretation, limitations, and future work.

2 Background

2.1 Classical random walks and hitting times

Let $G = (V, E)$ be a connected graph with $n = |V|$ vertices and let P be a row-stochastic transition matrix. A probability row vector π is stationary if $\pi P = \pi$. For a reversible chain, detailed balance holds: $\pi_x P_{xy} = \pi_y P_{yx}$ for all x, y.

Given a non-empty marked set $M \subset V$, the hitting time is the expected number of steps needed to reach M from the stationary distribution,

$$\mathrm{HT}(P, M) = \mathbb{E}_{X_0 \sim \pi} \left[\min\{t \geq 0 : X_t \in M\} \right]. \tag{1}$$

For computations we use the equivalent conditioned form

$$\mathrm{HT}(P, M) = \frac{1}{1 - p_M} \sum_{x \notin M} \pi_x h_x, \qquad h_x = 1 + \sum_y P_{xy} h_y, \quad x \notin M, \tag{2}$$

with $p_M = \sum_{x \in M} \pi_x$. The linear system in (2) is small for the graphs in this paper and can be solved directly. Standard examples are $\mathrm{HT} = \Theta(N)$ on the complete graph, $\mathrm{HT} = \Theta(N^2)$ on the cycle, and $\mathrm{HT} = \Theta(N \log N)$ on the two-dimensional grid; the hypercube behaves like $N \log N$ over the relevant range.

2.2 Szegedy walks and the detection–finding gap

A Szegedy walk lifts a Markov chain to a unitary operator on two vertex registers. With the convention used by AGJK,

$$V(P)|\bar{0}\rangle|x\rangle = \sum_y \sqrt{P_{xy}}|y, x\rangle, \qquad W(P) = V^\dagger(P)\,\mathrm{Shift}\,V(P)\,\mathrm{Ref}, \tag{3}$$

where $\mathrm{Shift}|x, y\rangle = |y, x\rangle$ and $\mathrm{Ref} = (2|\bar{0}\rangle\langle\bar{0}| - I) \otimes I$. The discriminant matrix is

$$D_{xy} = \sqrt{P_{xy} P_{yx}} = \langle\bar{0}|\langle x|W(P)|\bar{0}\rangle|y\rangle. \tag{4}$$

For a reversible chain, $D = \mathrm{diag}(\pi)^{1/2} P\, \mathrm{diag}(\pi)^{-1/2}$, so D and P share their spectrum and D is symmetric.

Szegedy's search algorithm detects a marked vertex in $O(\sqrt{\mathrm{HT}})$ steps. Detection is not the same as finding: measuring the final state need not return a marked vertex. The gap is especially visible when there are multiple marked vertices. This is the starting point of AGJK.

2.3 The AGJK algorithm and extended hitting time

AGJK introduce the absorbing walk P' and the interpolated walk $P(s) = (1 - s)P + sP'$, which stays at a marked vertex with probability s instead of moving. The corresponding quantum walk $W(s)$ uses an ancilla to record whether the current vertex is marked. AGJK define the extended hitting time HT^+ through the eigenvalues and eigenvectors of the discriminant matrix $D(s)$; if there is a single marked vertex, $\mathrm{HT}^+ = \mathrm{HT}$, but with multiple marked vertices HT^+ can be much larger.

Two algorithms are given. The simpler one runs $W(s)$ for $O(\sqrt{\mathrm{HT}})$ steps and is supported by a numerical conjecture (Conjecture 3 in the paper): for some s and $t = O(\sqrt{\mathrm{HT}})$, measuring the vertex register returns a marked vertex with constant probability. The proved algorithm combines the interpolated walk with quantum fast-forwarding and achieves $\widetilde{O}(\sqrt{\mathrm{HT}})$ time for any graph and any marked set (Ambainis et al., 2019). We implement and test the first algorithm numerically; we cite the second as the theoretical guarantee and do not re-prove it.

Paper	Model	Cognitive phenomenon	Role here
AGJK 2019	quantum walk search	marked-vertex search, HT^+	QIS core
Chen et al. 2022	DTQW	evidence accumulation, decoherence	measurement model
Huang–Busemeyer 2022	quantum walk	multi-alternative choice	decision mapping
Pavone–Faro 2025	CTQW on semantic graphs	insight, interference	semantic experiment
Pothos–Busemeyer 2022	quantum probability	context, order effects	motivation
Epping et al. 2023	open-system model	choice and response time	noise interpretation

Table 1: Literature map. The technical core is AGJK; the other papers supply the cognitive interpretation and the measurement/noise models.

2.4 Quantum cognition and quantum walks

Quantum cognition replaces classical probability with quantum probability in models of judgment and decision. Pothos and Busemeyer review the program and its explanations of order effects, conjunction and disjunction effects, and contextuality (Pothos and Busemeyer, 2022; Busemeyer and Bruza, 2024). The models are usually process-level: they do not assume that the brain implements a quantum circuit.

Quantum walks enter this literature in several ways. Chen, Ferro, and Sornette use discrete-time quantum walks as evidence accumulators and discuss how partial measurement produces decoherence (Chen et al., 2022). Huang and Busemeyer propose a quantum walk framework for multi-alternative decision making (Huang and Busemeyer, 2022). Pavone and Faro use continuous-time quantum walks on semantic graphs to model the emergence of insight through interference (Pavone and Faro, 2025). Epping et al. develop an open-system model of choice and response time (Epping et al., 2023), and Widdows et al. give quantum circuit components for cognitive decision-making (Widdows et al., 2023).

2.5 The gap

The two literatures meet at the level of metaphor but not at the level of algorithmic structure. Quantum-walk search papers analyze worst-case complexity and spectral quantities but do not ask what those quantities mean for a cognitive process. Quantum-cognition papers use walks as evidence accumulators but rarely exploit the marked-set structure, the extended hitting time, or the difference between detection and finding. This paper treats the marked-vertex search problem as a process model and uses its structural quantities—HT, HT^+, the spectral gap, and the success curve—as the bridge between the two literatures.

3 Theory

3.1 Reversible chains and the discriminant matrix

Let P be an ergodic reversible Markov chain with stationary distribution π. The discriminant matrix is

$$D_{xy} = \sqrt{P_{xy}P_{yx}}, \qquad D = \operatorname{diag}(\pi)^{1/2} P \operatorname{diag}(\pi)^{-1/2}. \tag{5}$$

For reversible P, D is symmetric and has the same spectrum as P; the vector $\sqrt{\pi}$ is a left eigenvector with eigenvalue 1. The Szegedy walk (3) satisfies

$$\langle\bar{0}|\langle x|W(P)|\bar{0}\rangle|y\rangle = D_{xy}. \tag{6}$$

The following standard spectral fact is used repeatedly. If $D = \sum_k \lambda_k |v_k\rangle\langle v_k|$, then the nontrivial eigenvalues of $W(P)$ are $e^{\pm i\theta_k}$ with $\lambda_k = \cos\theta_k$, and the corresponding eigenvectors lie in the span of $\{|\bar{0}\rangle|x\rangle\}_{x\in V}$ and its image under $V(P)$. This is the Szegedy spectral lemma; we use it as a statement and do not reprove it.

3.2 Classical and extended hitting times

For a marked set M with stationary mass $p_M = \sum_{x\in M}\pi_x$, the classical hitting time can be written in spectral form as

$$\mathrm{HT}(P,M) = \frac{1}{1-p_M}\sum_{k=1}^{n-|M|}\frac{|\langle v'_k|\sqrt{\pi_U}\rangle|^2}{1-\lambda'_k}, \tag{7}$$

where λ'_k are the non-unit eigenvalues of $D(1)$ and $\sqrt{\pi_U}$ is the subnormalized vector with entries $\sqrt{\pi_x}$ for $x \notin M$.

AGJK replace D by the discriminant matrix $D(s)$ of the interpolated walk and define

$$\mathrm{HT}(s) = \frac{1}{1-p_M}\sum_{k=1}^{n-1}\frac{|\langle v_k(s)|\sqrt{\pi_U}\rangle|^2}{1-\lambda_k(s)}, \qquad \mathrm{HT}^+ = \lim_{s\to 1}\mathrm{HT}(s), \tag{8}$$

where $D(s) = \sum_k \lambda_k(s)|v_k(s)\rangle\langle v_k(s)|$ and $\lambda_n(s) = 1$. For a single marked vertex, $\mathrm{HT}^+ = \mathrm{HT}$. For multiple marked vertices, (8) can be much larger than (7).

The canonical example is an $N \times N$ torus with marked set

$$M = M_1 \cup M_2, \qquad M_1 = \{(j_1d_1, j_2d_1) : 0 \le j_1, j_2 < k_1\}, \quad M_2 = \{(j_1d, j_2d) : 0 \le j_1, j_2 < N/d\}, \tag{9}$$

with $d_1 \mid d \mid N$ and the scale separation of AGJK. The set M_1 is a dense cluster and M_2 is a sparse periodic grid. For suitable parameters, $\mathrm{HT} = \widetilde{O}(d^2)$ while $\mathrm{HT}^+ = \Omega(N^2)$. The large gap comes from the periodic structure, not from the mere fact that $|M| > 1$.

3.3 Interpolated walk and the numerical AGJK protocol

Let $U = V \setminus M$. The absorbing and interpolated walks are

$$P' = \begin{pmatrix} P_{UU} & P_{UM} \\ 0 & I \end{pmatrix}, \qquad P(s) = (1-s)P + sP', \quad s \in [0,1). \tag{10}$$

The quantum version attaches a qubit that records whether the current vertex is marked. For $x \in U$, $V(P,s)$ acts as $V(P)$; for $x \in M$ it maps

$$|0\rangle|\bar{0}\rangle|x\rangle \mapsto \sqrt{1-s}|0\rangle V(P)|\bar{0}\rangle|x\rangle + \sqrt{s}|1\rangle|\bar{0}\rangle|x\rangle. \tag{11}$$

The interpolated walk operator is

$$W(s) = V^\dagger(P,s)\,\mathrm{Shift}'\,V(P,s)\,\mathrm{Ref}', \tag{12}$$

with the controlled shift and reflection of AGJK. Its top-left block is $D(s)$.

Algorithm 1 of AGJK prepares $|\bar{0}\rangle|\sqrt{\pi}\rangle$, applies $W(s)$ for t steps, and measures the marking ancilla. If the ancilla is 1, a position measurement returns a candidate marked vertex. The key numerical object is the lower bound

$$q_t(s) = \left\|\Pi_M T_t(D(s))|\sqrt{\pi}\rangle\right\|^2 \le \left\|(I \otimes \Pi_M)W^t(s)|\bar{0}\rangle|\sqrt{\pi}\rangle\right\|^2, \tag{13}$$

where T_t is the Chebyshev polynomial of the first kind. The identity behind (13) is

$$\Pi_0 W^t(s)\Pi_0 = |\bar{0}\rangle\langle\bar{0}| \otimes T_t(D(s)), \qquad \Pi_0 = |\bar{0}\rangle\langle\bar{0}| \otimes I. \tag{14}$$

In our numerical protocol we scan s over a fixed grid and choose the value that maximizes $\max_t q_t(s)$ over $t \le 3\sqrt{\mathrm{HT}} + 5$. We then run the exact operator $W(s)$ and record the actual success curve $p_M(t)$. This tests the numerical content of Conjecture 3 of AGJK without assuming it.

For completeness, the proved AGJK algorithm combines the interpolated walk with quantum fast-forwarding and runs in time

$$O\Big(S\sqrt{\log \mathrm{HT}} + \sqrt{\mathrm{HT}}(U + C)\sqrt{\log \mathrm{HT} \log\log \mathrm{HT}}\Big), \tag{15}$$

where S, U, and C are the costs of setup, update, and checking. We cite (15) as the theoretical guarantee; our experiments use the simpler Algorithm 1.

3.4 Partial measurement

The decoherence model is a position measurement with probability p before each unitary step. Let $\Pi_x = I \otimes |x\rangle\langle x|$ be the projector onto vertex x and let $R_x = |\bar{0}\rangle|x\rangle$ be the reset state. The one-step channel is

$$\rho_{t+1} = W(0)\left[(1-p)\rho_t + p\sum_x R_x\langle R_x|\rho_t|R_x\rangle\langle R_x|\right]W(0)^\dagger. \tag{16}$$

For $p = 0$ this is the pure quantum walk. For $p = 1$ the state is reset to a position basis state at every step; the position distribution then follows the classical Markov chain

$$Q_{xy} = \left\|\Pi_y W(0)|\bar{0}\rangle|x\rangle\right\|^2. \tag{17}$$

For $s = 0$, $Q = P$ and the measured walk is exactly the classical random walk.

Because (16) is linear in ρ, the average success probability can be computed without Monte Carlo. We precompute the pure states $|\phi_{x,a}\rangle = W(0)^a|\bar{0}\rangle|x\rangle$ and their position marginals and marked masses. A trajectory is summarized by the probability q_0 that no measurement has occurred and by the distribution $q_{x,a}$ over the last reset vertex x and age a. The update is a small linear recursion; the success probability is the corresponding mixture of the precomputed marked masses. The details are in appendix C.

3.5 Continuous-time walks on semantic graphs

For the semantic experiment we use the symmetric normalized Laplacian

$$L_{\mathrm{sym}} = I - D^{-1/2}AD^{-1/2} \tag{18}$$

and the random-walk generator $L_{\mathrm{rw}} = I - D^{-1}A$. The CTQW and CRW are

$$|\psi(t)\rangle = e^{-iL_{\mathrm{sym}}t}|c\rangle, \qquad \rho(t) = \rho(0)e^{-L_{\mathrm{rw}}t}, \tag{19}$$

where c is the cue vertex. The eigenvalues of L_{sym} and L_{rw} coincide, so comparing at $\tau = \lambda_1 t$ aligns the two relaxation scales. The target activation is $p_T(t) = |\langle T|\psi(t)\rangle|^2$ for CTQW and $p_T(t) = \rho_T(t)$ for CRW.

Walk object	Cognitive interpretation
Vertex $x \in V$	concept, evidence state, or option
Edge $(x, y) \in E$	semantic association or evidence transition
Marked set M	target memory or decision boundary
Initial state $\|\psi_0\rangle$	cue, question, or initial attention
Unitary $W(s)$ or e^{-iLt}	deliberation / evidence accumulation
Measurement	committing to an answer or decision
$p_M(t)$	target activation or success probability
$T_q = \arg\max_t p_M(t)$	response time at peak evidence
First threshold crossing	response time under a deadline
P_i for option sets M_i	choice probability
HT, HT$^+$	expected search difficulty, context difficulty
Spectral gap λ_1	deliberation / relaxation speed
Interpolation s	commitment or absorption strength
Measurement rate p	introspection or noise level

Table 2: Mapping between walk quantities and cognitive quantities. The mapping is process-level and mathematical; it is not a claim about the physical implementation of cognition.

4 Cognitive Model

4.1 From walk objects to cognitive quantities

The model treats a graph as a semantic or decision network and the walk as a deliberation process. Table 2 lists the correspondence. The mapping is deliberately minimal: every cognitive quantity is a statistic of the walk, so the model can be simulated and compared with a classical baseline without additional free parameters.

4.2 Modeling assumptions

We make five assumptions explicit.

1. **Process level.** The model describes the statistics of a search or decision process. It does not claim that the brain implements a quantum walk or a unitary circuit.

2. **Interference as multiple paths.** Quantum interference stands for the interaction of multiple evidence paths or semantic associations. The phase structure is a modeling device, not a physiological statement.

3. **Classical baseline.** The baseline is a Markov random walk on the same graph with the same initial distribution and the same marked set. This isolates the effect of interference from the effect of graph structure.

4. **Defined, not fitted, statistics.** Response time, choice probability, target activation, entropy, and radius are defined from the walk statistics. We do not fit behavioral data in this paper.

5. **Graph as prior.** The graph encodes the semantic or decision prior. Changing edges or weights changes the predicted context effects.

4.3 Response time and choice probability

For a single target, the activation curve is $p_M(t)$. We use two response-time definitions. The *peak time* is $T_q = \arg\max_{t \geq 1} p_M(t)$, the time at which the evidence for the target is strongest. The *threshold time* is the first t with $p_M(t) \geq \theta$, where θ is a decision threshold. The threshold definition is closer to a deadline experiment; the peak definition is closer to an unconstrained search. For decoherence we report both the peak success probability and the success probability at a fixed deadline.

For multiple option sets $M_1, \ldots, M_k$, the option probabilities at time t are

$$P_i(t) = \left\| (I \otimes \Pi_{M_i}) W^t(s) |\psi_0\rangle \right\|^2, \qquad P_i = \frac{P_i(T)}{\sum_j P_j(T)}, \tag{20}$$

where T is the decision time. The choice entropy $-\sum_i P_i \log P_i$ measures how concentrated the choice distribution is. A uniform distribution gives maximal entropy; a strong context effect gives low entropy and a large $\max_i P_i - 1/k$ bias.

4.4 Similarity and context effects

The decision experiments manipulate similarity by coupling two branches at matching depths with weight c. A large c means that the two options are strongly associated: a walker can move between them without returning to the hub. In a classical random walk this mainly redistributes probability among the coupled options. In a quantum walk the coupling changes the interference pattern, so the uncoupled option can gain disproportionately. This is a model-level analogue of similarity and attraction effects: the choice probability of an option depends on the presence and similarity of the other options, not only on its own value.

The semantic experiments manipulate distance in the same spirit. A cue and a target in the same community are close in the graph; a cross-community target requires a longer path. The CTQW and CRW are expected to differ most when the target is far from the cue, because the quantum walk maintains coherent amplitude across several paths while the classical walk relaxes locally.

4.5 Testable predictions

The model yields four predictions that can be tested numerically and, in principle, behaviorally.

1. For well-connected graphs and a single target, quantum search time scales as $\sqrt{\mathrm{HT}}$ and beats the classical baseline.

2. Independent targets are easier than clustered targets; a large HT^+ penalty requires a structured, not merely clustered, marked set.

3. Semantic distance changes activation timing; the quantum walk has a higher peak for far targets, while the classical walk has a higher early AUC for near targets.

4. Weak partial measurement can improve deadline success and shorten the time to half-maximum, but strong measurement destroys the quantum advantage.

These predictions are model-level. Confirming them numerically does not show that human cognition is quantum; it shows that the quantum-walk model makes distinct, falsifiable predictions relative to its classical counterpart.

ID	Graph / system	Marked set	Noise	Main metric
E1	K_N, C_N, grid, Q_d	single	none	scaling of T_q and HT
E2	grid, barbell	single / independent / clustered	none	$p_M(t)$, HT^+
E2b	adversarial torus	AGJK pattern	none	$\mathrm{HT}^+/\mathrm{HT}$
E3	toy semantic graph	single target	none	activation, AUC
E4	multi-branch graph	option terminals	none	choice prob., RT
E5	grid, toy semantic	single, clustered	partial measurement	success at deadline

Table 3: Experiment matrix. Each row corresponds to a processed result file in `code/results/processed/`.

5 Methods

5.1 Graph families and marked sets

We use four synthetic graph families for the scaling experiments: complete graphs K_N, cycles C_N, two-dimensional grids $L \times L$, and hypercubes Q_d with $N = 2^d$. The multiple-target experiments add barbell graphs (two cliques joined by a path) and $N \times N$ torus graphs with the marked-set pattern used in the AGJK counterexample. The semantic experiment uses a hand-built 19-vertex concept graph with animal, vehicle, and instrument communities, connected by the bridge concepts `bird` and `airplane`. Decision experiments use a hub with one path per option; matching depths of two branches can be coupled with a tunable weight.

For a graph $G = (V, E)$ we form the row-stochastic random-walk matrix

$$P_{xy} = \frac{A_{xy}}{\sum_z A_{xz}}, \tag{21}$$

where A is the adjacency matrix. Unless stated otherwise we use the lazy walk $P \mapsto (1-\eta)P + \eta I$ with $\eta = 0.1$ to avoid bipartite oscillations. A marked set $M \subset V$ is either a single target, an independent set of k vertices chosen by farthest-point sampling, a connected cluster grown by BFS, or a set of option terminals. Table 3 lists the experiment matrix.

5.2 Classical hitting time

For a marked set M, let h_x be the expected time to reach M from $x \notin M$, with $h_x = 0$ for $x \in M$. The hitting times solve the linear system

$$h_x = 1 + \sum_y P_{xy} h_y, \qquad x \notin M. \tag{22}$$

The classical hitting time is the stationary average conditioned on starting outside M,

$$\mathrm{HT}(P, M) = \frac{\sum_{x \notin M} \pi_x h_x}{1 - p_M}, \qquad p_M = \sum_{x \in M} \pi_x. \tag{23}$$

We solve (22) directly for all grids and barbell graphs; the synthetic families are small enough that dense linear algebra is inexpensive.

5.3 Szegedy and AGJK search protocol

The quantum walk uses the register convention of AGJK,

$$V(P)|\bar{0}\rangle|x\rangle = \sum_y \sqrt{P_{xy}}|y,x\rangle, \qquad W(P) = V^\dagger(P)\,\mathrm{Shift}\,V(P)\,\mathrm{Ref}, \tag{24}$$

with $\mathrm{Shift}|x,y\rangle = |y,x\rangle$ and $\mathrm{Ref} = (2|\bar{0}\rangle\langle\bar{0}| - I)\otimes I$. The discriminant matrix is $D_{xy} = \sqrt{P_{xy}P_{yx}}$.

For multiple or clustered targets we use the interpolated walk

$$P(s) = (1-s)P + sP', \qquad P' = \begin{pmatrix} P_{UU} & P_{UM} \\ 0 & I \end{pmatrix}, \tag{25}$$

and the corresponding walk operator $W(s)$ with an ancilla that records whether the current vertex is marked. The initial state is $|\bar{0}\rangle|\sqrt{\pi}\rangle$. For each graph and marked set we scan s over a fixed grid and select the value that maximizes the AGJK lower bound

$$q_t(s) = \left\|\Pi_M T_t(D(s))|\sqrt{\pi}\rangle\right\|^2, \qquad t \le 3\sqrt{\mathrm{HT}} + 5, \tag{26}$$

where T_t is the Chebyshev polynomial of the first kind. We then run the actual walk $W(s)$ and record

$$p_M(t) = \left\|(I\otimes\Pi_M)W^t(s)|\bar{0}\rangle|\sqrt{\pi}\rangle\right\|^2. \tag{27}$$

The quantum search time is $T_q = \arg\max_{t\ge1} p_M(t)$ and the peak success probability is $p_{\max} = \max_t p_M(t)$. The extended hitting time HT^+ is computed from the spectral definition

$$\mathrm{HT}(s) = \frac{1}{1-p_M}\sum_{k=1}^{n-1}\frac{|\langle v_k(s)|\sqrt{\pi_U}\rangle|^2}{1-\lambda_k(s)}, \qquad \mathrm{HT}^+ = \lim_{s\to1}\mathrm{HT}(s), \tag{28}$$

evaluated at $s = 1 - 10^{-6}$ as a numerical proxy. Equation (28) is used only for diagnostics; the search curves use the actual operator $W(s)$.

5.4 Semantic activation protocol

For the semantic experiment we use a continuous-time quantum walk (CTQW) and a continuous-time random walk (CRW) from the same cue vertex. Let

$$L_{\mathrm{sym}} = I - D^{-1/2}AD^{-1/2} \tag{29}$$

be the symmetric normalized Laplacian and $L_{\mathrm{rw}} = I - D^{-1}A$ the random-walk generator. The two models are

$$|\psi(t)\rangle = e^{-iL_{\mathrm{sym}}t}|c\rangle, \qquad \rho(t) = \rho(0)e^{-L_{\mathrm{rw}}t}, \tag{30}$$

with $\rho(0) = |c\rangle\langle c|$ at the cue c. Time is reported in units of the spectral gap, $\tau = \lambda_1 t$, so that both walks are compared at the same relaxation scale. The target activation is $p_T(t) = |\langle T|\psi(t)\rangle|^2$ for CTQW and $p_T(t) = \rho_T(t)$ for CRW. We record the peak activation, time of peak, area under the activation curve, entropy at the peak, half-peak coverage, and the expected graph distance from the cue at the peak.

5.5 Decision protocol

The decision graph has a central hub and one path per option. Each option i corresponds to a terminal set $M_i = \{\text{terminal}_i\}$. To control similarity, matching depths of branches i and j are joined by an edge of weight c; $c = 0$ gives independent branches and large c makes the two options highly similar. The quantum model runs the interpolated walk $W(s)$ with $s = 0.9$ and initial state $|\bar{0}\rangle|\text{hub}\rangle$. The classical model runs the same graph with all option terminals absorbing. For each model we compute the option probabilities $p_i(t)$ and the total decision probability $P_M(t) = \sum_i p_i(t)$. The decision time is the first t with $P_M(t) \geq 0.5$; if the threshold is never reached, the time of the maximum of $P_M(t)$ is used. Choice probabilities are the normalized option probabilities at the decision time.

5.6 Partial-measurement protocol

The decoherence experiment uses the plain Szegedy walk $W(0)$ from a cue vertex. Before each step, with probability p we measure the vertex register in the position basis and reset the state to $|\bar{0}\rangle|x\rangle$; with probability $1 - p$ the state evolves unitarily. Thus $p = 0$ is the pure quantum walk and $p = 1$ is exactly the classical random walk on the same graph. We compute the exact average over measurement trajectories by precomputing the position marginals and marked probabilities of the basis-state ensembles $W(0)^a|\bar{0}\rangle|x\rangle$, then propagating the age distribution of the most recent measurement. This avoids Monte Carlo noise and makes the sweep deterministic. The success probability is the probability that the walker is at a marked vertex after the step. We use $T_{\max} = 60$ and a deadline of $t = 20$.

5.7 Metrics and reproducibility

The main metrics are the classical hitting time HT, the extended hitting time HT^+, the quantum search time T_q, the peak success probability $p_{\max}$, the speedup HT/T_q, the target AUC, the activation entropy and radius, choice probabilities, decision time, and the success probability at a fixed deadline. Scaling exponents are fitted by ordinary least squares in log–log coordinates.

The implementation is in Python 3.12 with NumPy, SciPy, NetworkX, Matplotlib, PyYAML, and pytest. The environment is created with the Miniforge file `code/environment.yml` (environment `qwalk-cognition`). All random choices use explicit seeds; the configuration files are `code/configs/base.yaml` and `code/configs/experiments.yaml`. The test suite contains 47 tests covering unitarity, norm preservation, the classical limit, hitting-time identities, decision normalization, and the exact measurement ensemble. Every figure in this paper is regenerated by the commands listed in `code/README.md`.

6 Results

6.1 Classical versus quantum scaling (E1)

Figure 2 shows classical hitting times and quantum search times for the four synthetic families. The quantum time is the time of peak success probability, and the interpolation parameter is selected by (26). The fitted exponents are collected in table 4. Across all 16 instances the peak success probability is between 0.84 and 0.99, so the reported T_q values are not driven by low-probability fluctuations.

The cycle is the clearest illustration: its classical hitting time grows almost quadratically ($\alpha = 1.95$), while the quantum search time grows almost linearly ($\beta = 1.06$), giving a speedup that

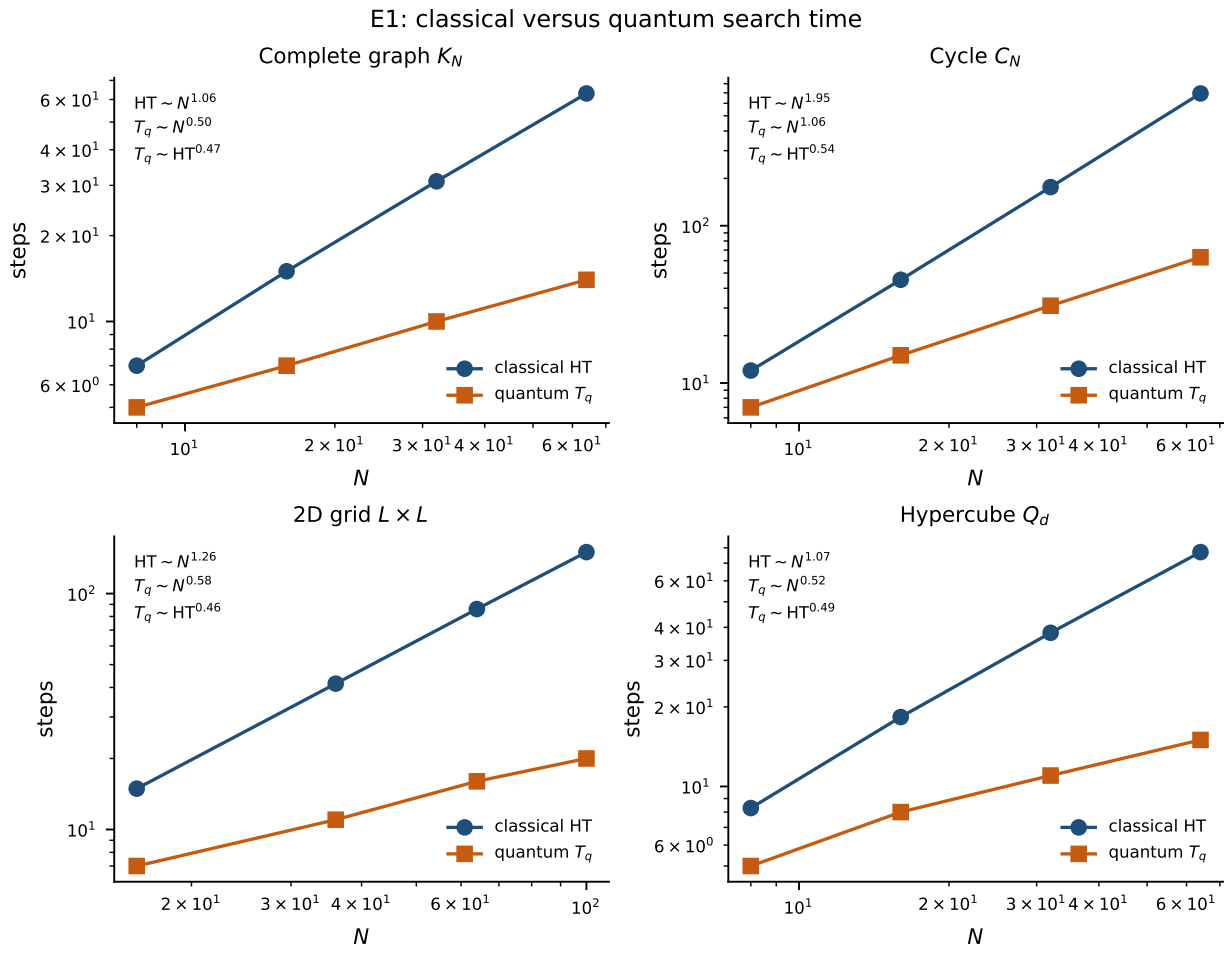


Figure 2: Classical hitting time and quantum search time versus graph size for complete graphs, cycles, 2D grids, and hypercubes. Both axes are logarithmic; lines are power-law fits. The quantum time tracks a square-root law in the classical hitting time rather than in the raw vertex count.

reaches 11× at $N = 64$. The complete graph and hypercube have HT $= \Theta(N \log N)$-like behavior over this size range, and the quantum exponents are again close to $1/2$. The grid shows a slightly larger exponent because the finite-size range does not yet fully separate N from $\log N$; the fit against HT gives $\gamma = 0.46$, consistent with the predicted square-root law.

Family	HT $\sim N^{\alpha}$	$T_q \sim N^{\beta}$	$T_q \sim \mathrm{HT}^{\gamma}$	max HT/T_q
Complete K_N	1.06	0.50	0.47	4.5
Cycle C_N	1.95	1.06	0.54	11.0
2D grid $L \times L$	1.26	0.58	0.46	7.5
Hypercube Q_d	1.07	0.52	0.49	5.1

Table 4: Fitted scaling exponents from E1. The exponent γ is the relevant comparison for the AGJK guarantee: it is close to 1/2 in all four families.

6.2 Single, independent, and clustered targets (E2)

Figure 3 compares success-probability curves for single, independent, and clustered marked sets. The independent configuration is consistently the easiest: on the 8×8 grid its mean HT is 31.2, versus 86.0 for a single target and 71.6 for a cluster of four. The same pattern is much stronger on the barbell graph, where independent targets reduce the hitting time from 2302 to 15.1 at $n_{\mathrm{each}} = 16$. Clustered targets behave like a single target: once the walker reaches the cluster it still has to find one of its vertices, so the hitting time remains close to the single-target value.

Figure 4 places classical hitting times against extended hitting times. For the grid and barbell cases the two quantities are close: the largest ratio is 1.52 for independent targets on the barbell, where several targets sit in different communities. The right panel shows the AGJK adversarial torus construction, where the ratio grows from 5.5 at $L = 18$ to 8.8 at $L = 36$. This confirms that a large HT^{+} penalty is a structural phenomenon: it requires a particular arrangement of marked vertices, not merely clustering. For $L = 18$ the interpolated walk still reaches $p_{\max} = 0.978$ at $t = 6$, so the Conjecture 3 numerical protocol is not slowed down by the large HT^{+} value.

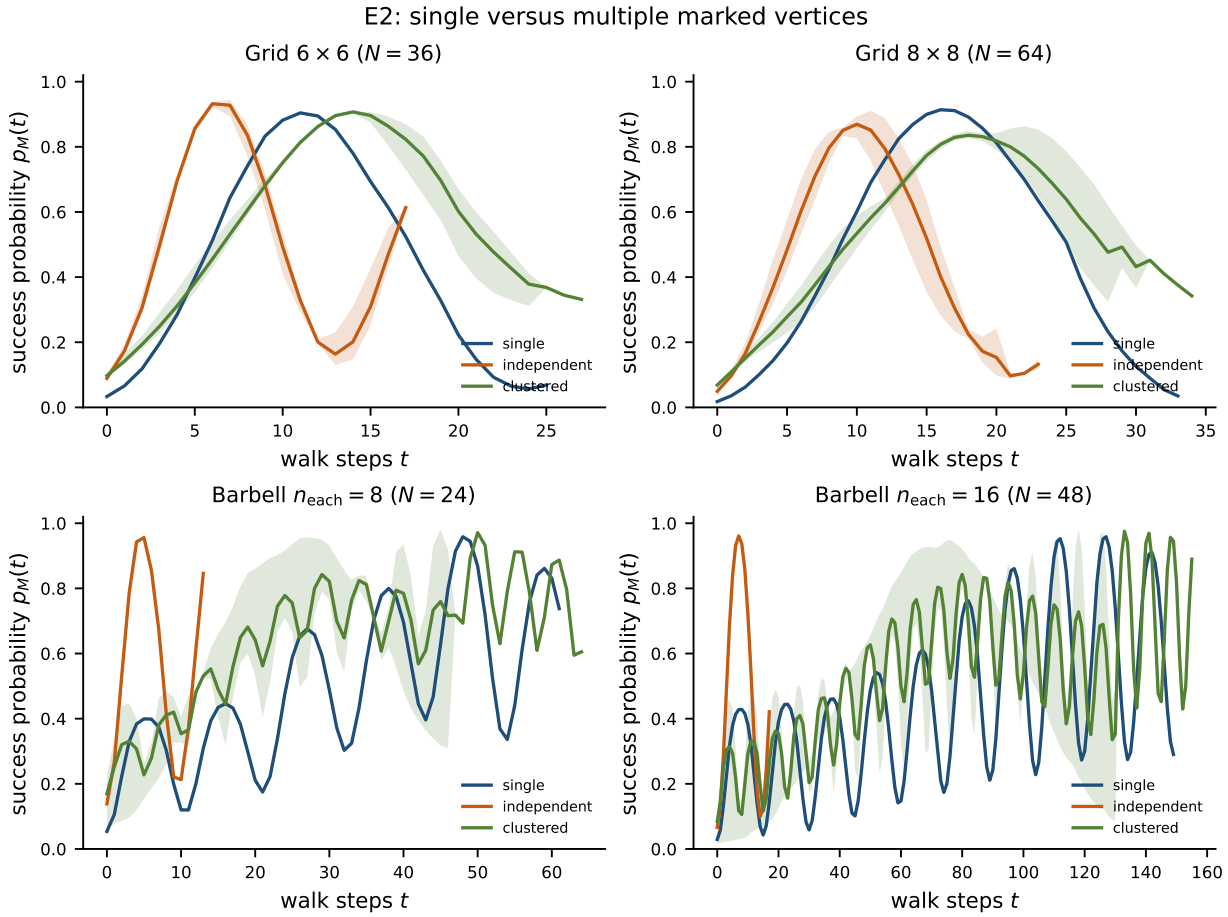


Figure 3: Success probability versus walk steps for single, independent, and clustered marked sets. Shaded bands show the min–max range over three random marked sets. Independent targets peak earlier and remain easier; clustered targets follow the single-target curve.

Graph	Marked set	mean HT	mean HT^+	$\mathrm{HT}^+/\mathrm{HT}$	mean T_q
Grid 8×8	single	86.0	86.0	1.00	16.0
Grid 8×8	independent	31.2	31.8	1.02	10.0
Grid 8×8	clustered	71.6	76.8	1.07	19.3
Barbell $n = 16$	single	2302.2	2302.1	1.00	127.0
Barbell $n = 16$	independent	15.1	22.9	1.52	7.0
Barbell $n = 16$	clustered	2230.7	2256.6	1.01	103.7

Table 5: Selected E2 results. Means are over three random marked sets for the independent and clustered configurations.

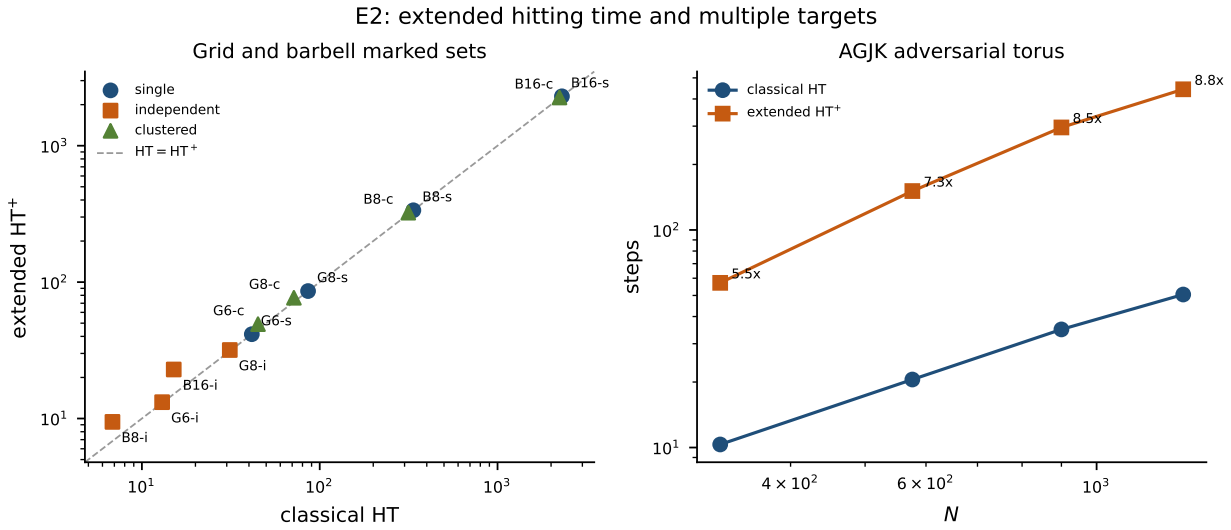


Figure 4: Left: classical hitting time versus extended hitting time for the grid and barbell marked sets; the dashed line is $\mathrm{HT} = \mathrm{HT}^{+}$. Right: the AGJK adversarial torus construction, where HT^{+} grows much faster than HT. Labels give the ratio $\mathrm{HT}^{+}/\mathrm{HT}$.

6.3 Semantic activation (E3)

Figure 5 shows the toy semantic graph and the target activation from the cue `animal` to two targets: `dog` (distance 1) and `violin` (distance 4). The spectral gap is $\lambda_1 = 0.0333$, giving a relaxation time of 30.0 steps. For the near target the CTQW peak is 0.237 at $\tau = 4.75$, compared with a CRW peak of 0.158 at $\tau = 0.13$. For the far target the gap is larger: CTQW reaches 0.247 at $\tau = 6.86$, while the CRW only reaches 0.046, essentially its stationary value.

The activation curves show a clear semantic-distance effect. For the near target the CRW has a higher area under the curve (0.580 versus 0.481) because it activates the neighbor immediately and then relaxes. For the far target the CTQW has a higher AUC (0.360 versus 0.324) and a much higher peak, while the CRW spreads almost uniformly. The expected radius at the CTQW peak is also larger (2.99 versus 2.36), indicating that the quantum walk maintains a broader, more coherent activation pattern instead of collapsing onto the local neighborhood.

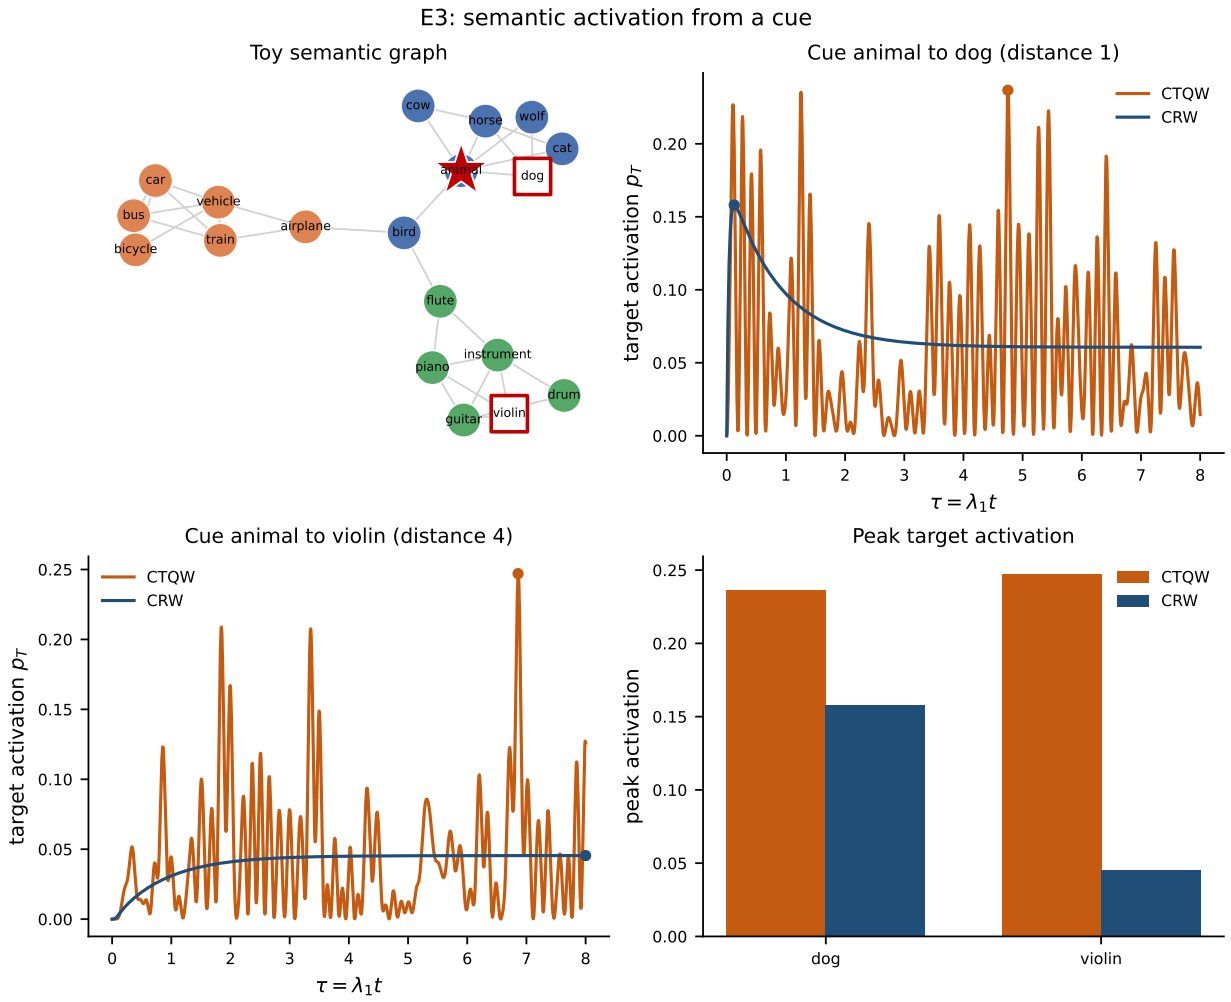


Figure 5: Toy semantic graph and target activation from the cue `animal`. Time is normalized by the spectral gap. The CTQW produces larger and later peaks than the CRW, especially for the cross-community target `violin`.

6.4 Multi-alternative choice under coupling (E4)

Figure 6 reports choice probabilities in three-option and four-option decision graphs as two branches are coupled with strength c. In the symmetric three-option case, increasing c from 0 to 2 moves the CTQW choice probability of the uncoupled option from 0.333 to 0.966, while the two coupled options fall to 0.017 each. The classical random walk remains close to uniform (0.394 and 0.303 at $c = 2$). The effect is not a small perturbation: coupling two options makes them behave as a joint alternative, and the quantum walker strongly prefers the alternative that is not part of the coupled pair.

The asymmetric case shows genuinely non-monotonic behavior. For the CTQW the choice probabilities of option A for $c = 0, 0.1, 0.25, 0.5, 0.75, 1, 2$ are $(0.145, 0.782, 0.094, 0.430, 0.955, 0.779, 0.902)$, and the decision times are $(5, 212, 186, 214, 89, 34, 17)$. The classical model is much less sensitive: its choice probabilities stay in the range 0.49–0.52, with decision times between 11 and 16. The symmetric four-option case behaves like the three-option case: the two uncoupled options gain (up

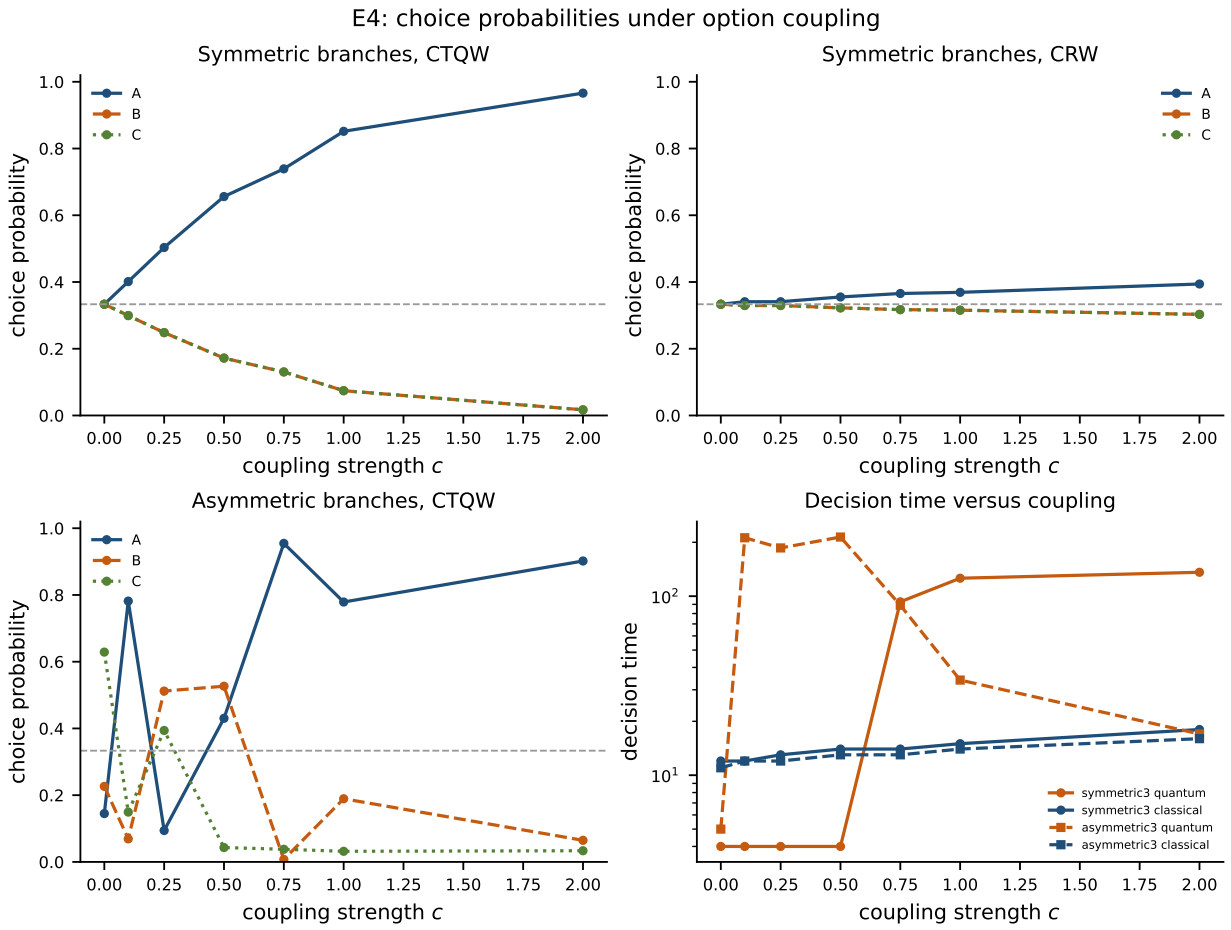


Figure 6: Choice probabilities and decision times under option coupling. Top: symmetric three-option graph for CTQW and CRW. Bottom left: asymmetric branches (3–4–5), where the CTQW choice probabilities are non-monotonic in c. Bottom right: decision time versus coupling.

to 0.492 each at $c = 2$), and the coupled pair is suppressed (down to 0.0077 each).

These results are a model-level analogue of similarity and context effects: when two alternatives become more similar, their individual choice probabilities are not simply averaged, and the remaining alternative can gain disproportionately. The effect is driven by interference in the quantum walk, not by a classical utility function.

6.5 Decoherence and partial measurement (E5)

Figure 7 shows the decoherence sweep. At a fixed deadline of $t = 20$, all four cases have the same optimal measurement rate, $p^* = 0.2$. The gain over the pure quantum walk is largest for the single-target grid (0.0025 to 0.0296, a factor of 11.9), and remains positive for the clustered grid (1.68×), the single-target semantic graph (1.31×), and the clustered semantic graph (1.21×). At the same time, the unrestricted peak success probability is highest at $p = 0$ in all four cases, so decoherence does not improve the best possible success probability.

The time-to-half-maximum reveals the speed side of the tradeoff. For the single-target semantic graph, weak decoherence reduces the half-maximum time from 53 steps to 8 steps, while strong

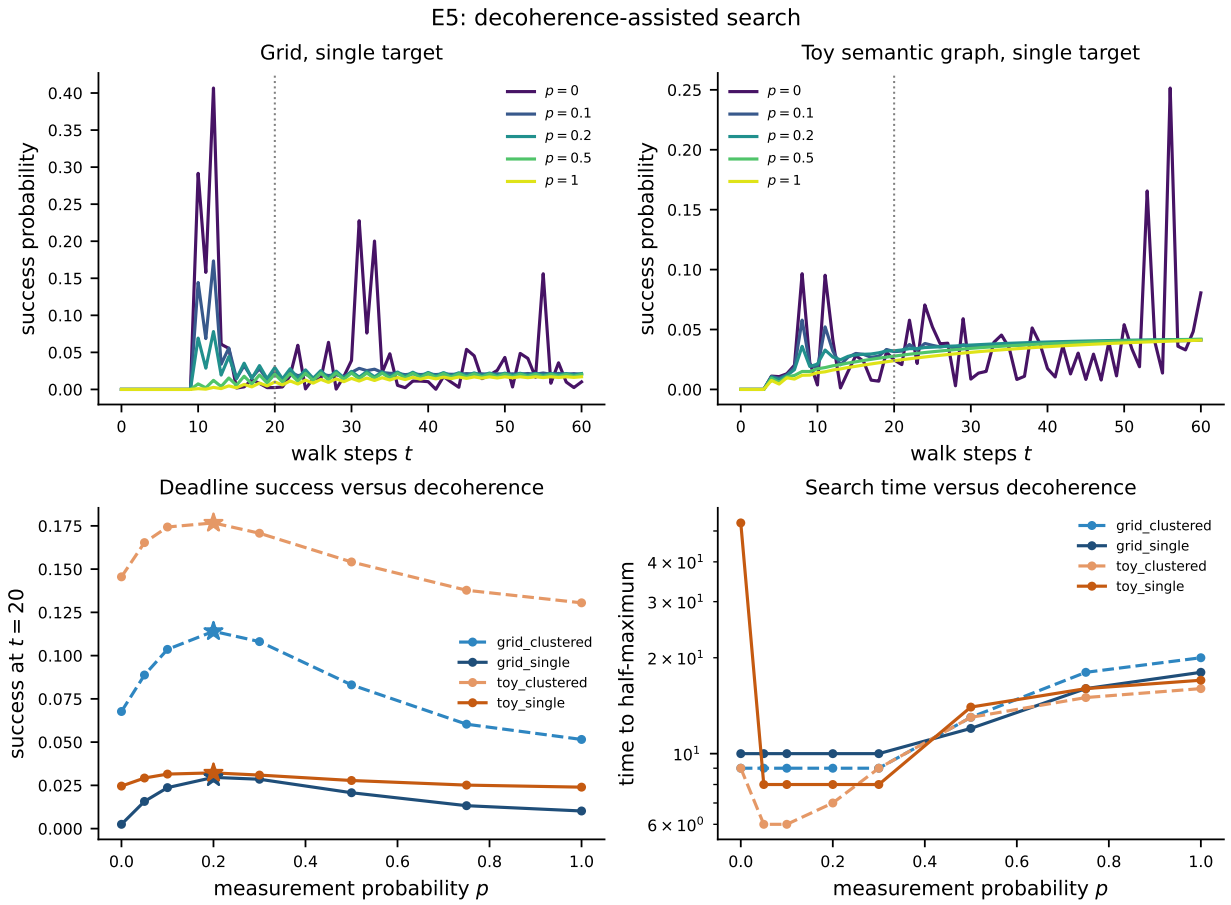


Figure 7: Partial-measurement decoherence sweep. Top: success probability versus time for selected measurement rates on the grid and semantic graph. Bottom left: success at the deadline $t = 20$ versus p; stars mark the optimum. Bottom right: time to half-maximum versus p.

measurement increases it again to 17 steps at $p = 1$. For the grid the half-maximum time is already short at $p = 0$ and changes little at weak p, but it grows at strong measurement. Thus weak decoherence can act as a useful "introspection" level: it suppresses destructive interference and accelerates the approach to a decision, but too much measurement destroys the quantum advantage.

Case	p^*	success at $t = 20$	ratio to $p = 0$	half-max time at $p = 0$	half-max time at p^*
Grid, single	0.2	0.0296	11.9	10	10
Grid, clustered	0.2	0.1140	1.68	9	9
Toy, single	0.2	0.0322	1.31	53	8
Toy, clustered	0.2	0.1767	1.21	9	7

Table 6: Optimal measurement rate at the deadline $t = 20$ and the associated speed–accuracy tradeoff.

6.6 Summary of empirical findings

The experiments support three conclusions. First, the quantum search time tracks $\sqrt{\mathrm{HT}}$ across the synthetic families, and independent targets are easier than clustered ones. Second, a large $\mathrm{HT}^+/\mathrm{HT}$ gap is a structural effect of the marked-set arrangement, not a generic consequence of having multiple targets. Third, the cognitive layer is sensitive to graph structure: semantic distance changes activation timing, option coupling produces strong context effects, and weak partial measurement can improve deadline performance and search speed even though it reduces the global peak success probability.

7 Discussion

7.1 What the experiments show

The scaling experiments reproduce the central AGJK phenomenon. Across four graph families the quantum search time fits $T_q \sim \mathrm{HT}^{0.46\text{–}0.54}$, close to the predicted square-root law, and the peak success probability is high enough that the fitted times are meaningful. The cycle gives the largest speedup because its classical hitting time is almost quadratic. The complete graph and hypercube have smaller finite-size speedups because their classical hitting times grow almost linearly over the simulated range.

The multiple-target experiments refine the picture. Independent targets are easier: spreading targets across the graph lowers both the classical hitting time and the quantum search time. Clustered targets behave like a single target, because reaching the cluster does not mean finding a particular vertex. The extended hitting time is close to the classical hitting time for these configurations. The dramatic $\mathrm{HT}^+ \gg \mathrm{HT}$ gap appears only in the periodic adversarial construction of AGJK. This is an important negative result for a naive reading of HT^+: multiple targets alone do not create the penalty. The penalty is a property of the arrangement and of the resulting spectral structure.

The semantic experiment shows a distance-dependent difference between CTQW and CRW. For a near target the classical walk activates the neighbor immediately and has a larger early AUC. For a far target the quantum walk maintains a larger and later peak, and its activation is spread over a larger radius. This is consistent with the intuition that coherent propagation can bridge communities, while classical diffusion relaxes to the stationary distribution. The result is a model-level prediction: a semantic-search experiment should see a larger peak but a later response for cross-community targets under the quantum model, and an earlier but lower activation under the classical model.

The decision experiment produces a strong context effect. Coupling two options makes them behave as a joint alternative, and the quantum walker shifts probability to the uncoupled option. In the symmetric case the effect is monotonic and large: the uncoupled option rises from 1/3 to 0.97

while the classical walk stays near 1/3. In the asymmetric case the choice probabilities and decision times are non-monotonic in the coupling strength. This is a model-level analogue of similarity and attraction effects. The classical model is much less sensitive because it lacks the interference that makes the coupled branches compete.

The decoherence experiment gives a nuanced answer to the noise question. At a fixed deadline, all four cases have an optimal measurement rate near $p = 0.2$. Weak measurement suppresses destructive interference and improves deadline success, by a factor of almost 12 in the single-target grid. At the same time, the unrestricted peak success probability is highest for the pure quantum walk. Weak measurement can also shorten the time to half-maximum dramatically, especially on the semantic graph. The model therefore predicts a speed–accuracy tradeoff rather than a single optimal noise level: the best amount of "introspection" depends on whether the observer is optimizing peak accuracy or decision speed.

7.2 Relation to the literature

The technical results are consistent with AGJK (Ambainis et al., 2019). Our numerical protocol tests the simpler interpolated-walk algorithm (Conjecture 3) and cites the fast-forwarding theorem as the proved guarantee. The adversarial torus reproduces the qualitative $\mathrm{HT}^{+} \gg \mathrm{HT}$ behavior and shows that the effect survives at modest sizes ($L = 18$–36). We do not claim a new theorem; the contribution is the numerical validation and the cognitive interpretation.

The semantic experiment is closest to Pavone and Faro (Pavone and Faro, 2025), but it uses a discrete marked-vertex search perspective rather than an insight task. The decision experiment is related to Huang and Busemeyer (Huang and Busemeyer, 2022) and to the evidence-accumulation literature (Chen et al., 2022). The decoherence model is a simplified version of the open-system ideas in Epping et al. (Epping et al., 2023): projective measurements at a fixed rate rather than a Lindblad master equation. The qualitative conclusion—that noise can help at intermediate levels and hurt at strong levels—is consistent with the broader quantum-cognition literature on decoherence-assisted processes.

7.3 Algorithmic versus biological claims

The project is deliberately agnostic about biology. Quantum walks are used as a mathematical process model, not as a claim that the brain is a quantum computer. The algorithmic advantage is a statement about the number of walk steps needed to reach a target under a particular graph and marked set. Whether a biological system can realize such a walk is a separate question. We keep the two claims separate throughout: the QIS results concern complexity and spectra, and the cognitive results concern the qualitative patterns that a process model can produce.

This separation also limits the interpretation of the context effects. The decision experiment shows that a quantum walk can produce non-monotonic choice probabilities under option coupling. It does not show that human similarity effects arise from quantum interference. The value of the model is that it generates distinct predictions: a classical random walk and a quantum walk on the same graph produce different choice distributions and different response times, so an experiment could in principle distinguish them.

7.4 Limitations

Several limitations should be stated clearly.

- **No human data.** The experiments are numerical. The cognitive quantities are model-defined, not fitted to behavioral data.
- **Synthetic graphs.** The semantic graph is hand-built and small (19 vertices). It illustrates the mechanism but does not capture the scale or noise of real semantic memory.
- **Interpolation selection.** The parameter s is chosen by a grid search on the AGJK lower bound. A different selection rule could change the constants in the search times, although the scaling exponents are stable.
- **Two walk formalisms.** The search experiments use discrete-time Szegedy walks, while the semantic experiment uses continuous-time walks. We compare them only within their own time normalizations and do not mix the two formalisms in a single fit.
- **Simplified noise.** The decoherence model uses projective position measurements and resets. It is not a full Lindblad or open-system model.
- **Finite sizes.** The graphs have at most a few hundred vertices. The asymptotic exponents are estimated from small ranges, especially for the grid and hypercube.

7.5 Future work

Four extensions would strengthen the project. First, replace the toy semantic graph with a ConceptNet subgraph and test whether the distance-dependent activation pattern survives in a real semantic network. Second, fit the decision model to behavioral choice and response-time data, using the coupling parameter as a free similarity parameter. Third, replace the projective measurement model with a Lindblad equation or a continuous dephasing model and compare the optimal noise level. Fourth, implement the decision graph on a small quantum device or circuit simulator to test whether the context effect survives realistic noise. A larger graph study could also test whether the non-monotonic choice probabilities in the asymmetric decision graph persist under different branch lengths and coupling patterns.

8 Conclusion

This project used the AGJK marked-vertex search algorithm as the QIS core of a process-level model of cognitive search. We implemented the Szegedy and interpolated walks, verified the $\widetilde{O}(\sqrt{\mathrm{HT}})$ scaling on four graph families, and compared single, independent, and clustered targets. The main structural finding is that independent targets are easier, while clustered targets behave like a single target; the large $\mathrm{HT}^{+} \gg \mathrm{HT}$ gap requires the adversarial periodic arrangement of AGJK rather than clustering alone. On the cognitive side, a toy semantic graph produced distance-dependent activation patterns, a multi-branch decision graph produced strong and non-monotonic context effects under option coupling, and a partial-measurement model produced an optimal measurement rate at a fixed deadline with a clear speed–accuracy tradeoff.

The broader lesson is that the algorithmic structure of quantum walk search is a useful source of cognitive hypotheses. The marked-set structure, the extended hitting time, and the success curve are not merely complexity quantities; they define a process model with testable predictions. At the same time, the project is deliberately cautious: the results are numerical, model-level, and do not claim a biological quantum advantage. The next step is to test the same predictions on larger semantic networks and, eventually, on behavioral choice and response-time data.

AI Tool Usage

An AI coding assistant was used for code scaffolding, experiment automation, and LaTeX drafting. All technical derivations, experiment design, numerical results, and interpretations were checked by the authors. The AI tool was not used to generate research results or citations without verification.

A Hitting-time derivations

Complete graph. For K_N with one marked vertex, every unmarked vertex is equivalent. Let h be the expected hitting time from an unmarked vertex. Then

$$h = 1 + \frac{N-2}{N-1}h, \tag{31}$$

so $h = N - 1$. Since the stationary distribution is uniform and we condition on starting outside the marked set, $\mathrm{HT} = N - 1$.

Cycle. For C_N with marked vertex 0, the expected hitting time from distance i is $h_i = i(N-i)$. It satisfies the recurrence

$$h_i = 1 + \frac{1}{2}h_{i-1} + \frac{1}{2}h_{i+1}, \qquad h_0 = h_N = 0, \tag{32}$$

whose solution is the quadratic above. Averaging over $i = 1, \ldots, N-1$ gives

$$\mathrm{HT} = \frac{1}{N-1}\sum_{i=1}^{N-1} i(N-i) = \frac{N(N+1)}{6}. \tag{33}$$

These identities are used as sanity checks in `tests/test_classical.py`.

B Discriminant matrix and spectral identity

For a reversible chain, detailed balance gives $\pi_x P_{xy} = \pi_y P_{yx}$, hence

$$D_{xy} = \sqrt{P_{xy}P_{yx}} = \frac{\sqrt{\pi_x}P_{xy}}{\sqrt{\pi_y}}, \tag{34}$$

so $D = \mathrm{diag}(\pi)^{1/2} P\, \mathrm{diag}(\pi)^{-1/2}$. The matrices D and P are similar, so they have the same eigenvalues, and D is symmetric. The identity

$$\langle\bar{0}|\langle x|W(P)|\bar{0}\rangle|y\rangle = \sqrt{P_{xy}P_{yx}} = D_{xy} \tag{35}$$

follows directly from the definitions of $V(P)$, Shift, and Ref. For the interpolated walk, the same computation gives the top-left block $D(s)$, which is the starting point for the Chebyshev bound (13).

Graph	Marked set	mean HT	mean HT^+	$\mathrm{HT}^+/\mathrm{HT}$	mean T_q
Grid 6×6	single	41.5	41.5	1.00	11.0
Grid 6×6	independent	13.0	13.2	1.01	6.7
Grid 6×6	clustered	45.0	49.2	1.09	14.0
Grid 8×8	single	86.0	86.0	1.00	16.0
Grid 8×8	independent	31.2	31.8	1.02	10.0
Grid 8×8	clustered	71.6	76.8	1.07	19.3
Barbell $n = 8$	single	336.2	336.2	1.00	48.0
Barbell $n = 8$	independent	6.8	9.4	1.38	5.0
Barbell $n = 8$	clustered	314.8	322.9	1.04	39.0
Barbell $n = 16$	single	2302.2	2302.1	1.00	127.0
Barbell $n = 16$	independent	15.1	22.9	1.52	7.0
Barbell $n = 16$	clustered	2230.7	2256.6	1.01	103.7

Table 7: Full E2 table. Means are over three random marked sets for the independent and clustered configurations.

C Measurement ensemble recursion

Let $|\phi_{x,a}\rangle = W(0)^a|\bar{0}\rangle|x\rangle$ for $a = 0, \ldots, T_{\max}$, with position marginal $\mu_{x,a}(y) = \|\Pi_y|\phi_{x,a}\rangle\|^2$ and marked mass $s_{x,a} = \sum_{y \in M} \mu_{x,a}(y)$. Let $\mu_a^{(0)}$ and $s_a^{(0)}$ be the corresponding quantities for the initial state. A trajectory is described by $q_0(t)$, the probability that no measurement has occurred by time t, and by $q_{x,a}(t)$, the probability that the last reset was at vertex x and the state has evolved for a steps since that reset.

At one step,

$$q_0(t+1) = (1-p)q_0(t), \tag{36}$$

$$q_{x,a+1}(t+1) = (1-p)q_{x,a}(t), \tag{37}$$

$$q_{y,1}(t+1) = p\, q_0(t)\mu_t^{(0)}(y) + p \sum_{x,a} q_{x,a}(t)\mu_{x,a}(y). \tag{38}$$

The success probability is

$$p_M(t+1) = q_0(t+1)s_{t+1}^{(0)} + \sum_{x,a} q_{x,a}(t+1)s_{x,a}. \tag{39}$$

This recursion is exact for the channel (16). It requires the basis ensembles once per graph and marked set, after which a sweep over p is inexpensive. The implementation is in `src/noise.py`.

D Additional results

Figure 8 shows the classical-to-quantum speedup for E1. The cycle has the largest speedup because its classical hitting time grows almost quadratically; the complete graph has the smallest because its classical hitting time grows only linearly.

E Reproducibility

The environment is created with

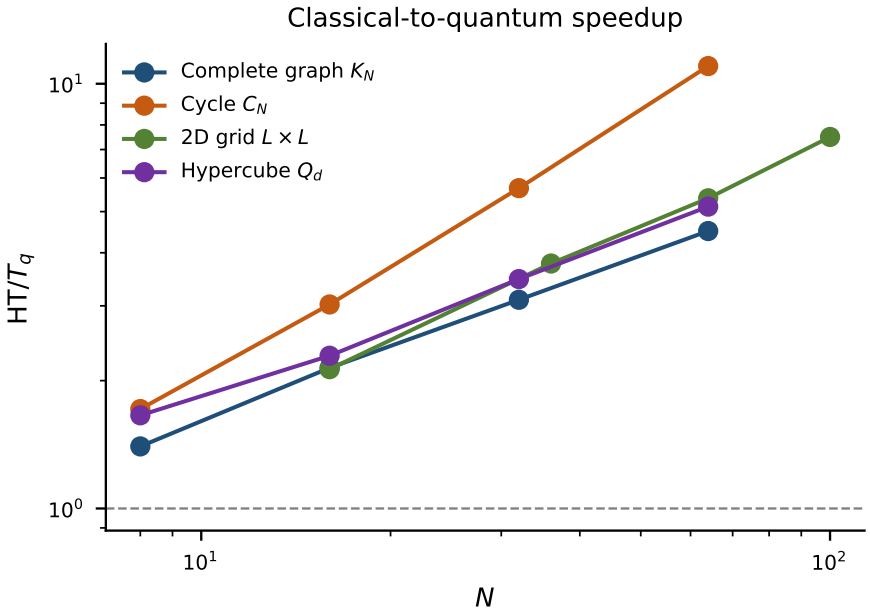


Figure 8: Classical-to-quantum speedup HT/T_q versus graph size for E1.

```
mamba env create -f environment.yml
mamba activate qwalk-cognition
```

from `Final/code`. The experiments are run with

```
python -m src.experiments --only E1
python -m src.experiments --only E2
python -m src.experiments --only E3
python -m src.experiments --only E4
python -m src.experiments --only E5
```

and write the CSV/JSON files in `results/` and the figures in `figures/`. The test suite is run with `python -m pytest -q` and currently contains 47 passing tests. The configuration files are `configs/base.yaml` and `configs/experiments.yaml`; all random choices use explicit seeds.

References

Andris Ambainis, András Gilyén, Stacey Jeffery, and Martins Kokainis. Quadratic speedup for finding marked vertices by quantum walks. *arXiv preprint arXiv:1903.07493*, 2019.

Jerome R. Busemeyer and Peter D. Bruza. *Quantum Models of Cognition and Decision*. Cambridge University Press, 2 edition, 2024. doi: 10.1017/9781009205351.

Ming Chen, Gustavo M. Ferro, and Didier Sornette. On the use of discrete-time quantum walks in decision theory. *PLoS ONE*, 17(8):e0273551, 2022. doi: 10.1371/journal.pone.0273551.

Gijs P. Epping, Peter D. Kvam, Timothy J. Pleskac, and Jerome R. Busemeyer. Open system model of choice and response time. *Journal of Choice Modelling*, 49:100453, 2023. doi: 10.1016/j.jocm.2023.100453.

Jerry Huang and Jerome R. Busemeyer. A quantum walk framework for multialternative decision making. *Proceedings of the 44th Annual Meeting of the Cognitive Science Society*, 2022. doi: 10.31234/osf.io/hmv4a.

Alessandro Pavone and Simone Faro. Modeling the emergence of insight via quantum interference on semantic graphs. *Mathematics*, 13(19):3171, 2025. doi: 10.3390/math13193171.

Emmanuel M. Pothos and Jerome R. Busemeyer. Quantum cognition. *Annual Review of Psychology*, 73:749–778, 2022. doi: 10.1146/annurev-psych-033020-123501.

Dominic Widdows, Jyoti Rani, and Emmanuel M. Pothos. Quantum circuit components for cognitive decision-making. *Entropy*, 25(4):548, 2023. doi: 10.3390/e25040548.